

Group SBS, Sheet 03, Exercise 02

November 19, 2021

Contributors

Sreyas Sony (soney@uni-potsdam.de), Binoy Chacko (chacko@uni-potsdam.de),
Dinesh Kumar (kumar@uni-potsdam.de) Sanika Nair (nair@uni-potsdam.de)

Solution

1. when X_1, \dots, X_n are independent and identical distributed random variables following the Bernoulli distribution.

θ is the unknown parameter associated with the probability of heads. The pmf of the Bernoulli distribution is thus given by,

$$f(x_i; \theta) = \theta^{x_i} (1 - \theta)^{1-x_i}$$

The Likelihood function is given by,

$$\begin{aligned} L(\theta) &= \prod_{i=1}^n f(x_i; \theta) \\ &= \theta^{x_1} (1 - \theta)^{1-x_1} \times \theta^{x_2} (1 - \theta)^{1-x_2} \times \dots \times \theta^{x_n} (1 - \theta)^{1-x_n} \\ &= \theta^{\sum_{i=1}^n x_i} (1 - \theta)^{n - \sum_{i=1}^n x_i} \end{aligned}$$

Taking the log of $L(\theta)$,

$$\log(L(\theta)) = \log(\theta) \sum_{i=1}^n x_i + \left(n - \sum_{i=1}^n x_i \right) \log(1 - \theta)$$

Differentiating w.r.t. θ ,

$$\frac{\partial}{\partial \theta} \log(L(\theta)) = \frac{\sum_{i=1}^n x_i}{\theta} - \frac{n - \sum_{i=1}^n x_i}{1 - \theta}$$

Equating $\frac{\partial}{\partial \theta} \log(L(\theta))$ to 0,

$$\begin{aligned} \frac{\sum_{i=1}^n x_i}{\theta} - \frac{n - \sum_{i=1}^n x_i}{1 - \theta} &= 0 \\ \sum_{i=1}^n x_i - \theta \sum_{i=1}^n x_i - \theta n + \theta \sum_{i=1}^n x_i &= 0 \\ \Rightarrow \hat{\theta} &= \frac{1}{n} \sum_{i=1}^n x_i \end{aligned}$$

To check if $\hat{\theta}$ maximizes the likelihood, taking the 2nd derivative of $\log(L(\theta))$ w.r.t. θ ,

$$\frac{\partial^2}{\partial \theta^2} \log(L(\theta)) = -\frac{\sum_{i=1}^n x_i}{\theta^2} - \frac{n - \sum_{i=1}^n x_i}{(1 - \theta)^2}$$

Substituting the value of $\hat{\theta}$ in θ in the above equation,

$$\begin{aligned} &= -\frac{n^2}{\sum_{i=1}^n x_i} - \frac{n^2}{n - \sum_{i=1}^n x_i} \\ &= \frac{-n^3 + n^2 \sum_{i=1}^n x_i - n^2 \sum_{i=1}^n x_i}{n \sum_{i=1}^n x_i - (\sum_{i=1}^n x_i)^2} \\ &= \frac{-n^3}{n \sum_{i=1}^n x_i - (\sum_{i=1}^n x_i)^2} < 0 \end{aligned}$$

Thus, $\hat{\theta} = \frac{1}{n} \sum_{i=1}^n x_i$ maximizes the likelihood.

2. when X_1, \dots, X_n are independent and identical distributed random variables following the Binomial distribution.

θ is the unknown parameter associated with the probability of heads. The pmf of the Binomial distribution is thus given by,

$$\begin{aligned} f(x_i; \theta) &= \binom{n}{x_i} \theta^{x_i} (1 - \theta)^{n - x_i} \\ L(\theta) &= \prod_{i=1}^n f(x_i; \theta) \\ &= \prod_{i=1}^n \binom{n}{x_i} \times \left[\theta^{\sum_{i=1}^n x_i} (1 - \theta)^{n^2 - \sum_{i=1}^n x_i} \right] \end{aligned}$$

Taking the log of $L(\theta)$,

$$\log(L(\theta)) = \log\left(\prod_{i=1}^n \binom{n}{x_i}\right) + \log(\theta) \sum_{i=1}^n x_i + (n^2 - \sum_{i=1}^n x_i) \log(1 - \theta)$$

Differentiating w.r.t. θ ,

$$\frac{\partial}{\partial \theta} \log(L(\theta)) = 0 + \frac{1}{\theta} \sum_{i=1}^n x_i - \frac{1}{1 - \theta} \left(n^2 - \sum_{i=1}^n x_i \right)$$

Equating $\frac{\partial}{\partial \theta} \log(L(\theta))$ to 0,

$$\frac{\sum_{i=1}^n x_i - \theta \sum_{i=1}^n x_i - \theta n^2 + \theta \sum_{i=1}^n x_i}{\theta(1 - \theta)} = 0$$

$$\sum_{i=1}^n x_i - \theta n^2 = 0$$

$$\Rightarrow \hat{\theta} = \frac{1}{n^2} \sum_{i=1}^n x_i$$

To check if $\hat{\theta}$ maximizes the likelihood, taking the 2nd derivative of $\log(L(\theta))$ w.r.t. θ ,

$$\frac{\partial^2}{\partial \theta^2} \log(L(\theta)) = -\frac{\sum_{i=1}^n x_i}{\theta^2} - \frac{n^2 - \sum_{i=1}^n x_i}{(1 - \theta)^2}$$

Substituting the value of $\hat{\theta}$ in θ in the above equation,

$$\begin{aligned} &= -\frac{n^4}{\sum_{i=1}^n x_i} - \frac{n^4}{n^2 - \sum_{i=1}^n x_i} \\ &= \frac{-n^6 + n^4 \sum_{i=1}^n x_i - n^4 \sum_{i=1}^n x_i}{n^2 \sum_{i=1}^n x_i - (\sum_{i=1}^n x_i)^2} \\ &= \frac{-n^6}{n^2 \sum_{i=1}^n x_i - (\sum_{i=1}^n x_i)^2} < 0 \end{aligned}$$

Thus, $\hat{\theta} = \frac{1}{n^2} \sum_{i=1}^n x_i$ maximizes the likelihood.

Additional: When $X = X_1 + \dots + X_n$ is Binomial

In this case, each X_i is an iid Bernoulli trial, then X is the total sum of successes from each trial.

$$f(x; \theta) = \binom{n}{x} \theta^x (1 - \theta)^{n-x}$$

The Likelihood function is given by,

$$\begin{aligned} L(X|\theta) &= f(x; \theta) \\ &= \binom{n}{x} \theta^x (1 - \theta)^{n-x} \end{aligned}$$

Taking the log of $L(\theta)$,

$$\log(L(\theta)) = \log \left(\binom{n}{x} \right) + x \log(\theta) + (n - x) \log(1 - \theta)$$

Differentiating w.r.t. θ ,

$$\frac{\partial}{\partial \theta} \log(L(\theta)) = \frac{x}{\theta} - \frac{n - x}{1 - \theta}$$

Equating $\frac{\partial}{\partial \theta} \log(L(\theta))$ to 0,

$$\begin{aligned} \frac{x}{\theta} - \frac{n - x}{1 - \theta} &= 0 \\ x - x\theta - n\theta + x\theta &= 0 \\ \Rightarrow \hat{\theta} &= \frac{x}{n} \\ &= \frac{1}{n} \sum_{i=1}^n x_i \text{ because } X = X_1 + X_2 + \dots + X_n \end{aligned}$$

As proved in the 1st part, taking the second derivative proves that $\hat{\theta}$ maximizes the likelihood.

We can see that, the MLE based on n independent Bernoulli random variables and the MLE based on a single binomial random variable are the same, since binomial is the result of n independent Bernoulli trials anyway. In general, whenever we have repeated independent Bernoulli trials with the same probability of success θ for each trial, the MLE will always be the sample proportion of successes. This is true regardless of whether we know the outcomes of the individual trials X_1, X_2, \dots, X_n or just the total number of successes for all trials $X = X_1 + X_2 + \dots + X_n$.

Python Implemenation

November 19, 2021

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```


```
[2]: # Loading the data
data = pd.read_csv('sampleset.txt', header=None)[0].tolist()
print(data)
```

```
[1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1,
0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0,
0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0,
1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1]
```

```
[3]: sample_average = np.mean(data)
print(f'The sample average is {sample_average}')
```

The sample average is 0.43

```
[4]: # Plotting the distribution of the given dataset
height = np.histogram(data, bins=2)
plt.figure(figsize=(4,4))
plt.bar(x=['0', '1'], height=height[0], width=0.3)
plt.ylabel("Frequency")
plt.suptitle("Distribution of the dataset")
plt.tight_layout()
```



0.1 Parameter estimation using Bernoulli distribution

0.1.1 Theroretical method

$$\text{MLE} = \frac{1}{n} \sum_{i=1}^n x_i$$

This is equal to the sample average


```
[5]: theta_estimate_Ber = np.mean(data)
      print(f'The estimate for p is {theta_estimate_Ber}')
```

The estimate for p is 0.43

0.1.2 Graphical method

```
[6]: def likelihood_Ber(counts, n=100):
      """
      counts: total number of success
      n: number of trials
      """
      p = np.round(np.linspace(0,1,100), 2)
      #likelihood function
      L = p**counts * (1-p)**(n-counts)
      return p, L
```

```
[7]: # Total number of success in the given dataset
# number of trials
n = 100
success = np.sum(data)
x, y = likelihood_Ber(counts=success, n=100)
# getting the index value where likelihood is maximum
idx = np.argmax(y)
plt.plot(x, y)
plt.axvline(x[idx], ls='--', c='orange', label= f'MLE estimator = {x[idx]}')
plt.xlabel("p")
plt.ylabel("Likelihood")
plt.title("Probability vs Likelihood function")
plt.legend()
plt.tight_layout()
```



0.2 Parameter estimation using Binomial distribution

0.2.1 Theroretical method

Considering number of success is the sum of n i.i.d. Bernoulli trials, which is then a single Binomial random variable.

$$X = X_1 + X_2 + X_3 + \dots + X_n$$

$$\text{MLE} = \frac{1}{n} \sum_{i=1}^n x_i$$

the MLE based on n i.i.d. Bernoulli random variables and the MLE based on a single binomial random variable are the same, since the binomial is the result of n independent Bernoulli trials.

0.2.2 Parameter estimation using Binomial distribution


```
[8]: theta_estimate_Binom = np.mean(data)
     print(f'The estimate for p is {theta_estimate_Binom}')
```

The estimate for p is 0.43

Since the likelihood function is same for b

```
[9]: from math import factorial
     def likelihood_Binom(counts, n=100):
         """
         counts: total number of success
         n: number of trials
         """
         p = np.round(np.linspace(0,1,100), 2)
         #likelihood function
         nCk = factorial(n)/(factorial(counts)*factorial(n-counts))
         L = nCk*p**counts * (1-p)**(n-counts)
         return p, L
```

```
[10]: # Total number of success in the given dataset
      # number of trials
      n = 100
      success = np.sum(data)
      x, y = likelihood_Binom(counts=success, n=n)
      # getting the index value where likelihood is maximum
      idx = np.argmax(y)
      plt.plot(x, y)
      plt.axvline(x[idx], ls='--', c='orange', label= f'MLE estimator = {x[idx]}')
      plt.xlabel("p")
      plt.ylabel("Likelihood")
      plt.title("Probability vs Likelihood function")
      plt.legend()
      plt.tight_layout()
```

The MLE based on n i.i.d. Bernoulli random variables and the MLE based on a single binomial random variable are the same, since the binomial is the result of n i.i.d. Bernoulli trials. The difference between the likelihood function graph for Bernoulli and Binomial is only in the value of likelihood, which is due to the constant term $(n!/(k!(n-k)!))$ in the Binomial distribution.

[]: